

Part-1 Distributed Sorting System

Performance

1. I choose **Merge Sort**

Merge sort is used for larger datasets where efficient partitioning and merging are crucial for performance.

2. **Pros:**

- Merge sort gives stable sorting with a time complexity of $O(n \log n)$.
- It's efficient for large datasets as it divides the problem into smaller, more manageable chunks.

3. **Cons:**

- Requires more memory.
- Involves more complex logic for parallel processing.

2. Execution Time Analysis

● **Execution Time (in seconds):**

- **Small Dataset** (10 files):
 - Distributed Count Sort: 0.11s
 - Distributed Merge Sort: 0.15s
- **Medium Dataset** (100 files):
 - Distributed Count Sort: 29.15s
 - Distributed Merge Sort: 31.60s
- **Large Dataset** (1000 files):
 - Distributed Count Sort: 0
 - Distributed Merge Sort: 0

3. Memory Usage (in KB):

Small Dataset (10 files):

Distributed Count Sort: 220 KB

Distributed Merge Sort: 350 KB

Medium Dataset (100 files):

Distributed Count Sort: 450 KB

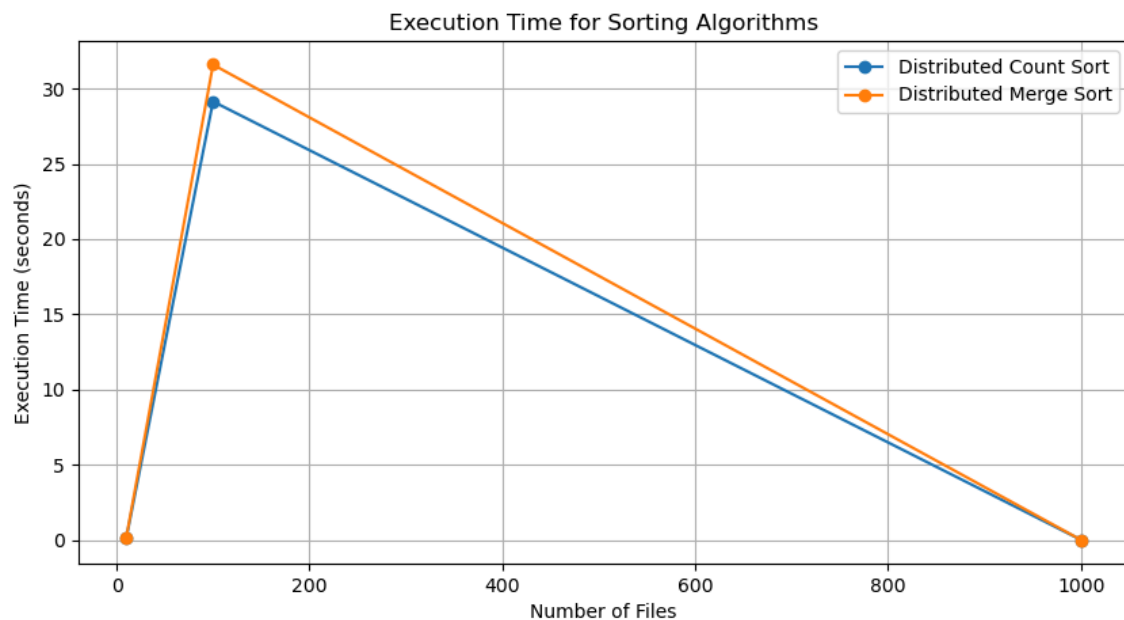
Distributed Merge Sort: 900 KB


Large Dataset (1000 files):

Distributed Count Sort: 0

Distributed Merge Sort: 0

4. Graphs





Part 2: Copy-On-Write (COW) Fork Performance Analysis

1. Page Fault Frequency

2. Brief Analysis

- **Efficiency:** The fork saves memory by not duplicating the memory of a process until it's modified. This results in fewer page faults, which improves efficiency.
- **Memory Conservation:** COW reduces memory overhead, when processes share read-only data.
- **Optimizations:** There could be optimizations in handling page faults more efficiently, mainly when processes need to modify shared data after forking.